



EXPLAINER

How AI search works, for implementers

A deeper explainer for the people who build and maintain websites: how assistants retrieve and cite content, and the technical choices that affect whether you appear.

PREPARED BY THE DIGITAL TOURISM THINK TANK

For implementers. This assumes you build or maintain sites. It explains how an assistant gets from a question to a cited answer, and the concrete technical choices that move you toward being the source it quotes.

01 From question to answer

At a useful level of detail, an assistant does four things. It retrieves candidate content it can access, often via live search and its own index. It ranks that content for relevance and clarity. It synthesises an answer from the strongest, most consistent sources. And it attributes, naming or linking the sources it leaned on. Each stage is a place you can win or lose: be retrievable, be clear, be consistent, be citable.

02 Overviews and assistants differ

Google's AI Overviews are grounded heavily in live web results and favour content that already performs in search. The standalone assistants blend trained knowledge with live retrieval, and lean on a smaller set of trusted, well structured sources. The practical implication is the same: clean, retrievable, well structured pages serve both, but do not assume one engine behaves like another, and test each.

03 Crawler access

If an assistant's crawler cannot fetch a page, that page cannot be cited. The named agents to account for include GPTBot, Google-Extended, ClaudeBot and PerplexityBot. Decide access deliberately in robots.txt: allowing them is what lets your content be used in answers; blocking them removes you from that surface. Treat it as a visibility decision, not a default.

04 Structured data, done right

Schema removes ambiguity about what a fact is. Implement the types that match your content, your organisation or place, events, offers, and validate with Google's Rich Results Test. Two rules matter. The schema must match the visible content exactly. And note that Google retired FAQ rich results in 2026 and assistants give FAQ markup no special weight, so use schema to describe what is genuinely on the page, never as a visibility device.

05 Content structure

Write answer-first. Open the page, and each section, with the direct answer, then elaborate. Use headings phrased as the questions users ask, keep sections self contained so they survive being extracted out of context, and put the load-bearing facts in text rather than in images or PDFs. Specificity helps: named places, times, prices and distances are exactly what an assistant lifts.

06 Entity and consistency

Assistants resolve you to an entity and trust facts that agree across sources. Keep name, address and core facts identical across your site, Google Business Profile and the major platforms. Use sameAs links to connect your profiles, and for genuinely notable entities, a correct, neutral, well sourced Wikidata entry helps the systems identify you. Contradictions are what get you dropped or misdescribed.

07 Measurement

Define a priority question set and run it across the main assistants on a schedule, recording where you appear and what is cited. Watch your server logs for the AI crawlers to confirm access and frequency. In analytics, track referral traffic from assistant domains. Trend over time is the signal; a single reading is noise.

THE LINE NOT TO CROSS

Cloaking, hidden text, prompt injection in page content, and fabricated reviews are against platform rules, detectable, and damaging when caught. None of them is GEO. The whole method is to make genuinely good, well structured information easy to retrieve and trust.

Prepared by the Digital Tourism Think Tank as part of the AI Discoverability programme. Crawler names and platform behaviours change, so verify current details before implementation.